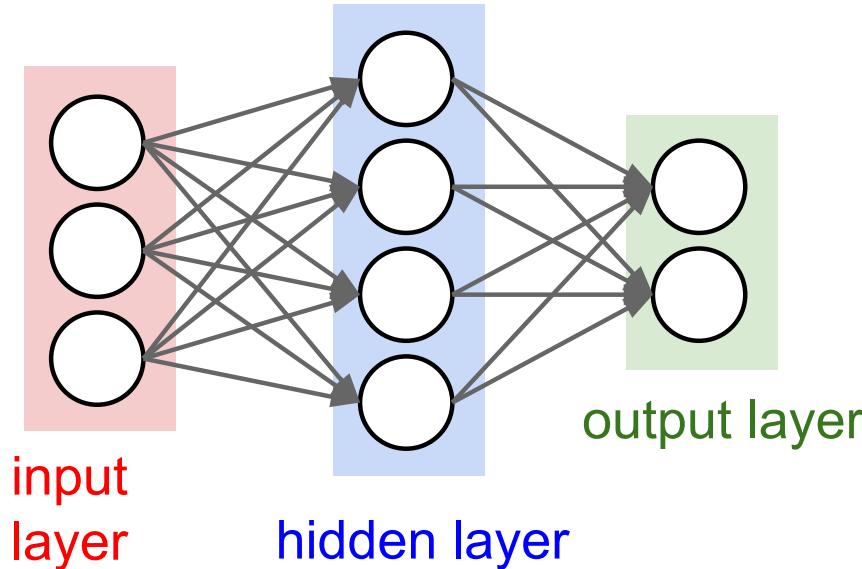
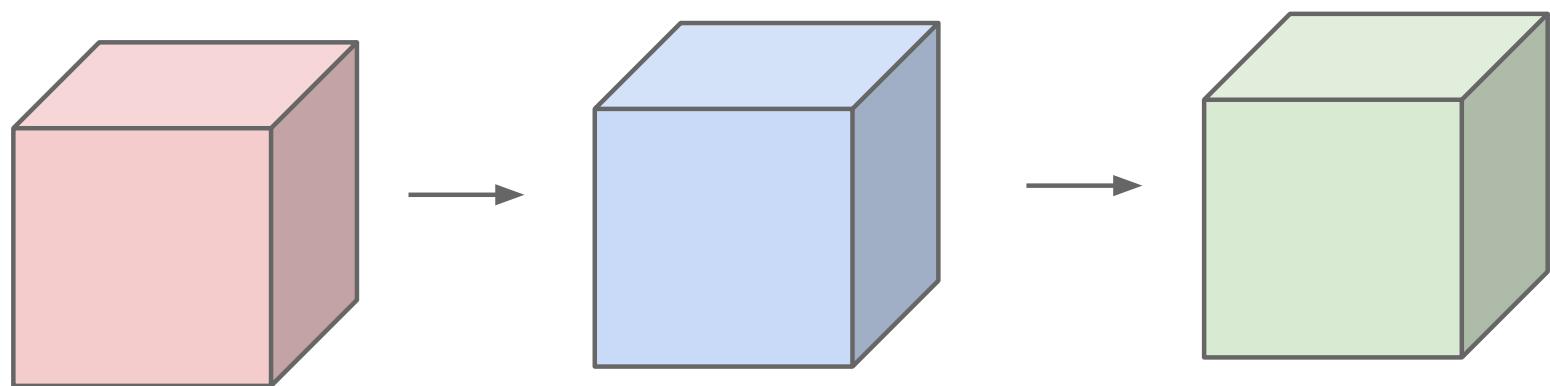




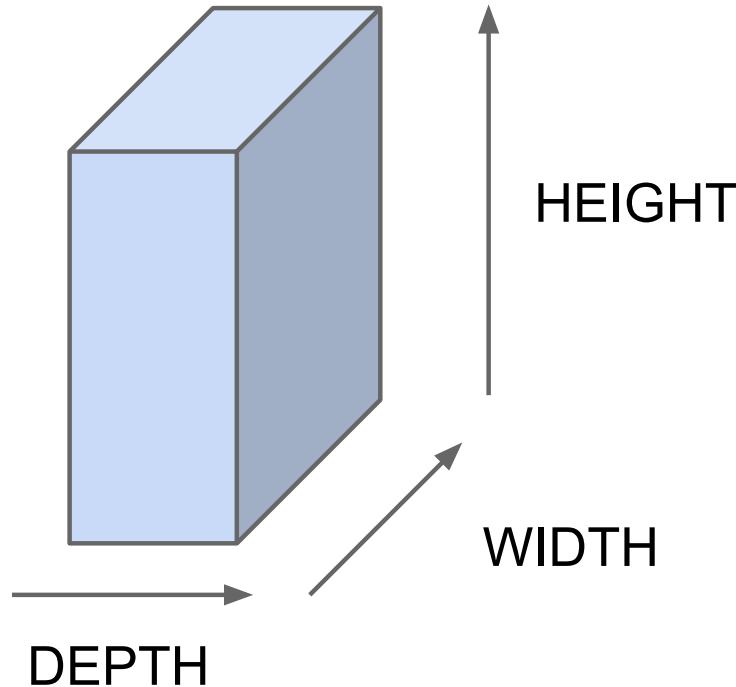
before:



now:

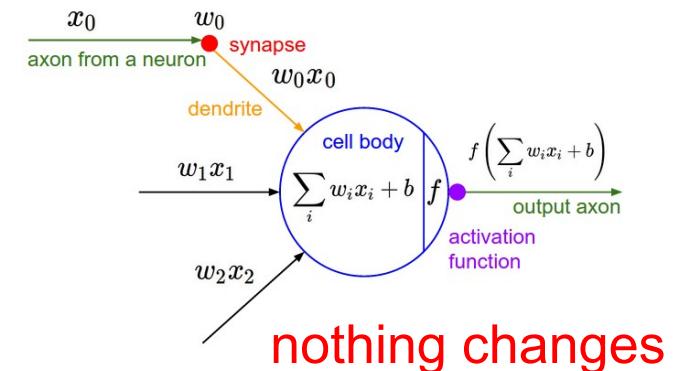
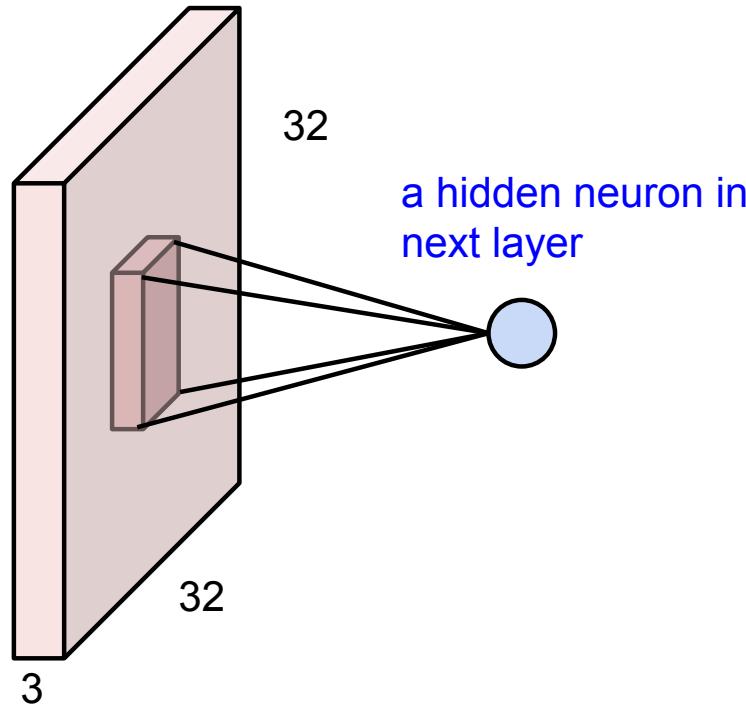


All Neural Net
activations
arranged in **3**
dimensions:



For example, a CIFAR-10 image is a $32 \times 32 \times 3$ volume
32 width, 32 height, 3 depth (RGB channels)

Convolutional Neural Networks
are just Neural Networks BUT:
1. Local connectivity



Convolutional Neural Networks
are just Neural Networks BUT:
1. Local connectivity

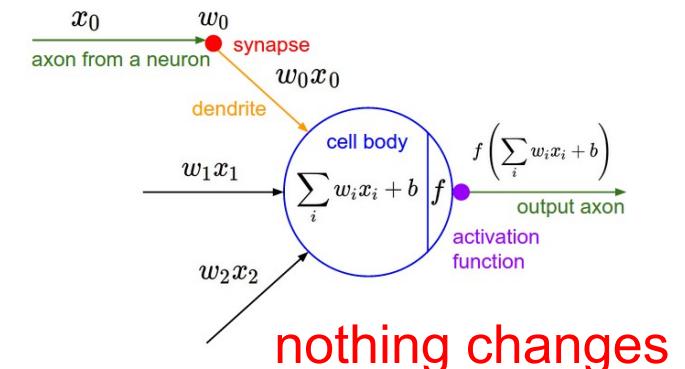
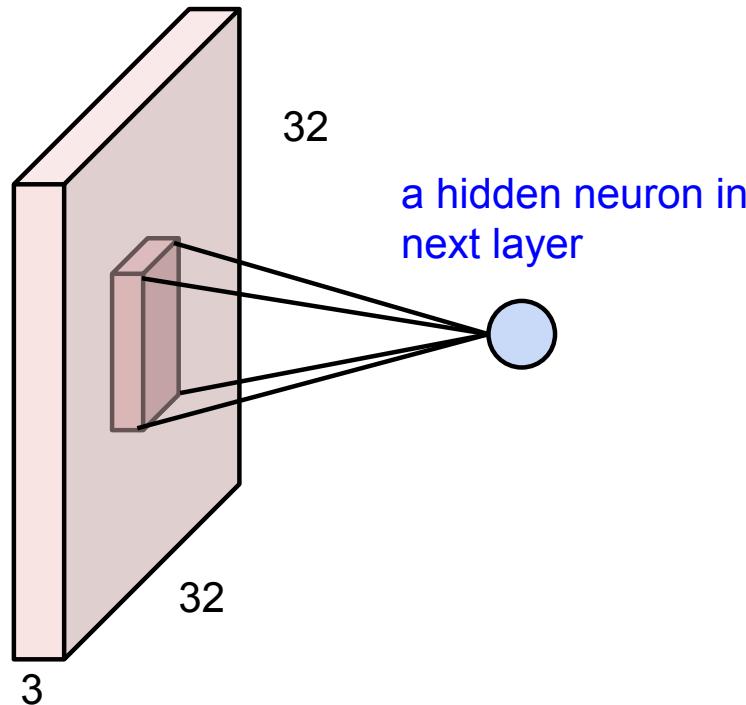


image: 32x32x3 volume

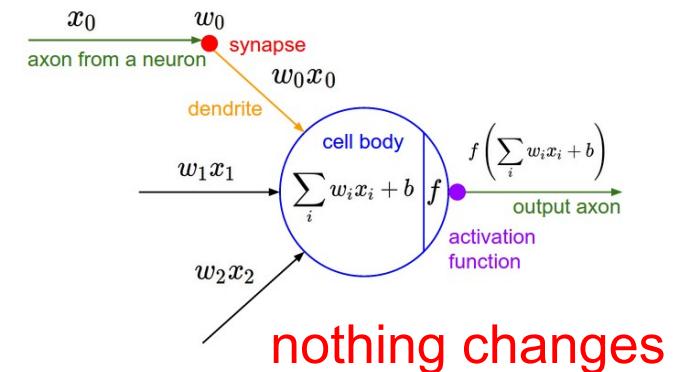
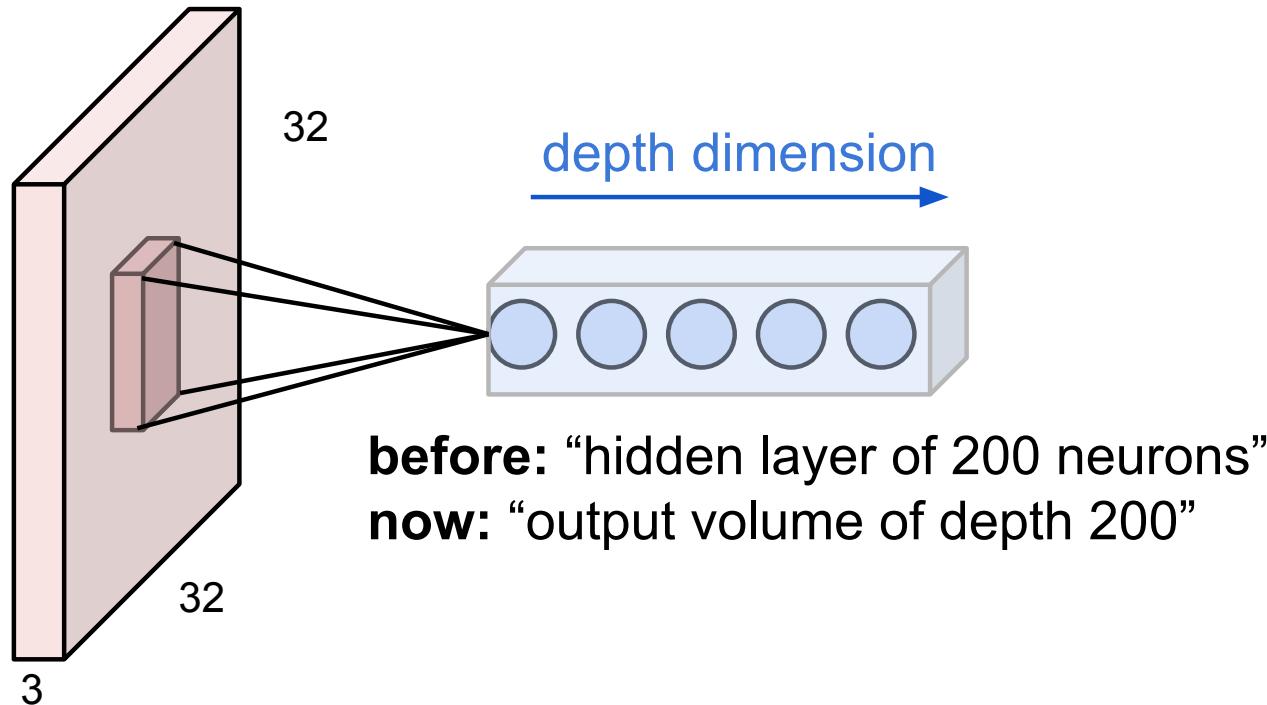
before: full connectivity: 32x32x3 weights

now: one neuron will connect to, e.g. 5x5x3 chunk and only have 5x5x3 weights.

note that connectivity is:

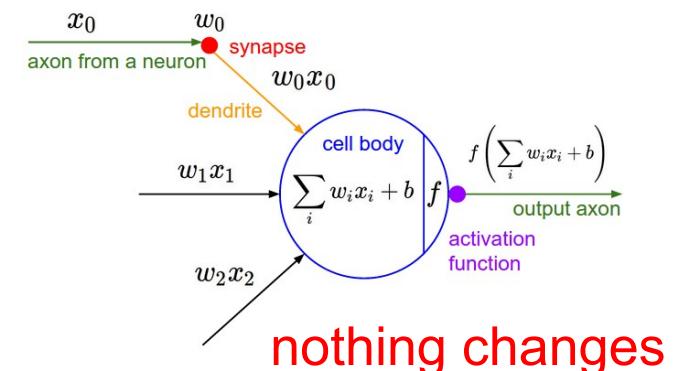
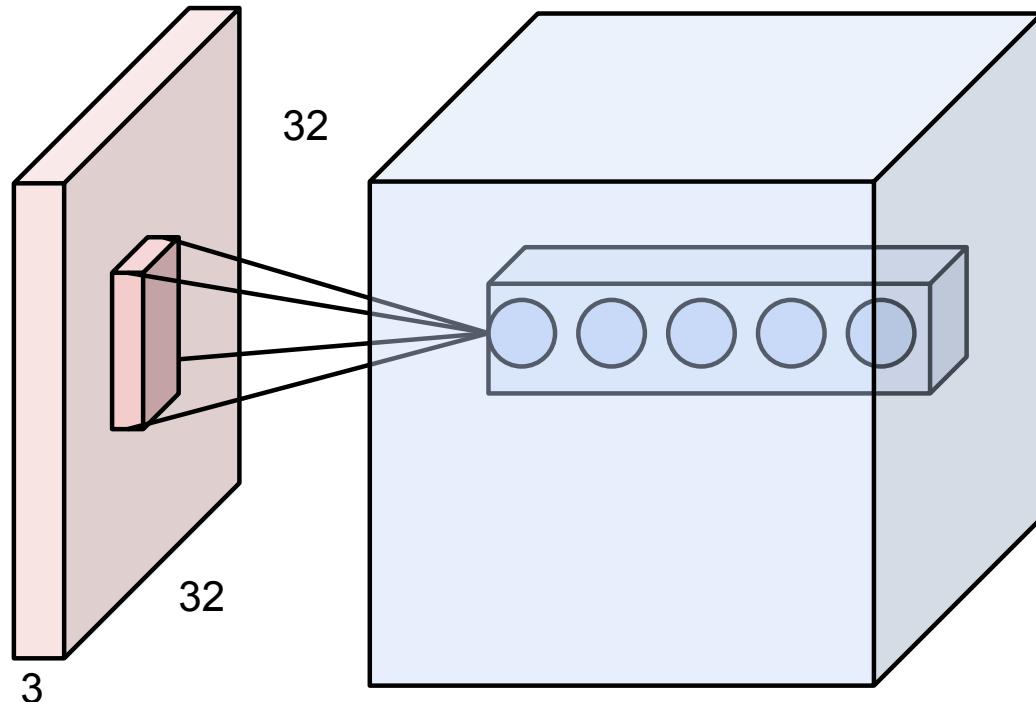
- local in space (5x5 inside 32x32)
- but full in depth (all 3 depth channels)

Convolutional Neural Networks
are just Neural Networks BUT:
1. Local connectivity



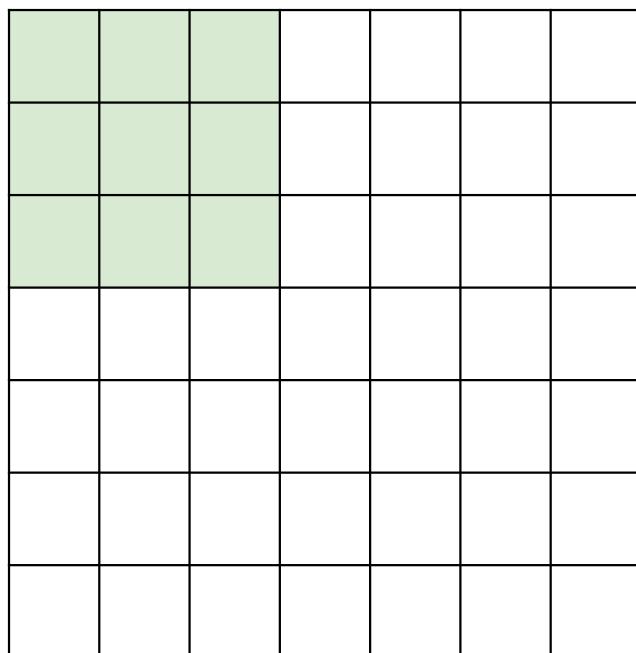
Multiple neurons all
looking at the same
region of the input
volume, stacked
along depth.

Convolutional Neural Networks
are just Neural Networks BUT:
1. Local connectivity



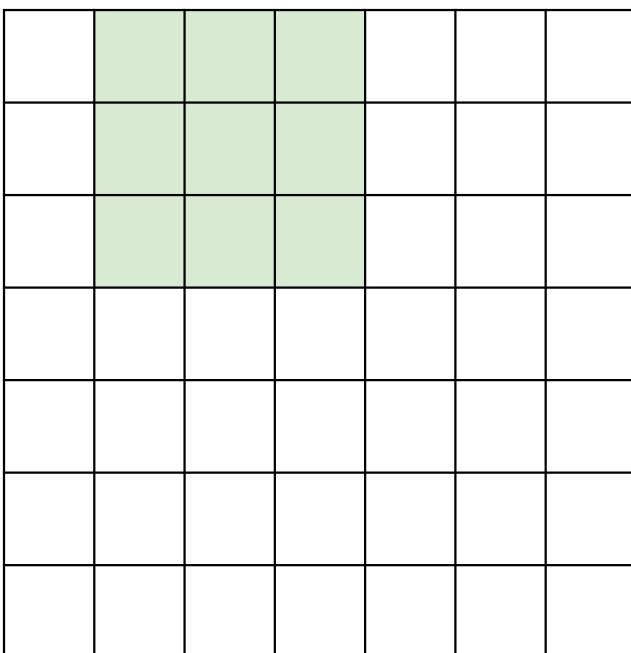
These form a single
[1 x 1 x depth]
“depth column” in the
output volume

Replicate this column of hidden neurons across space, with some **stride**.



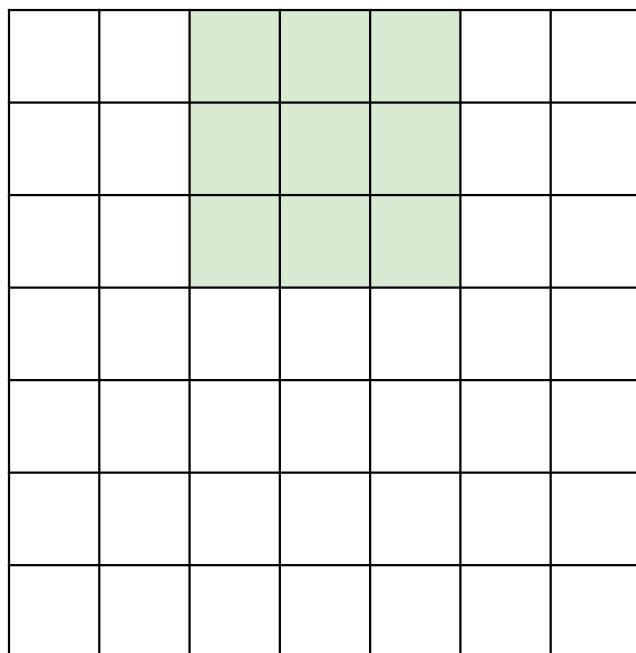
7x7 input
assume 3x3 connectivity, stride 1

Replicate this column of hidden neurons across space, with some **stride**.



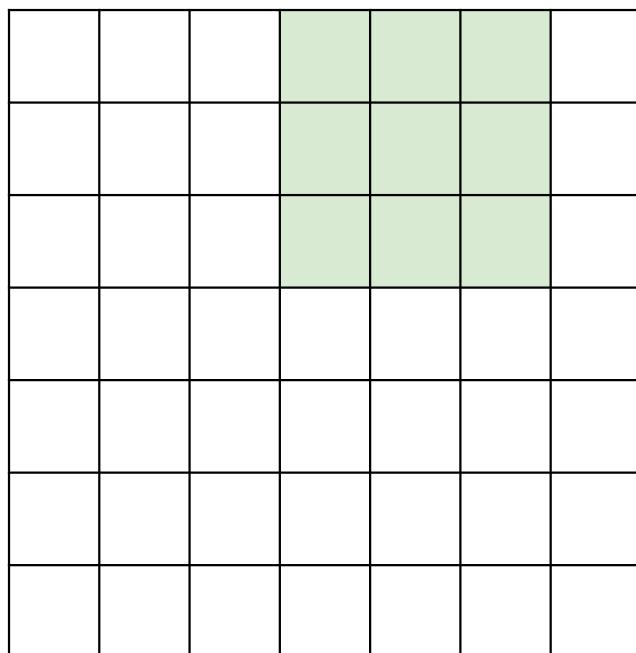
7x7 input
assume 3x3 connectivity, stride 1

Replicate this column of hidden neurons across space, with some **stride**.



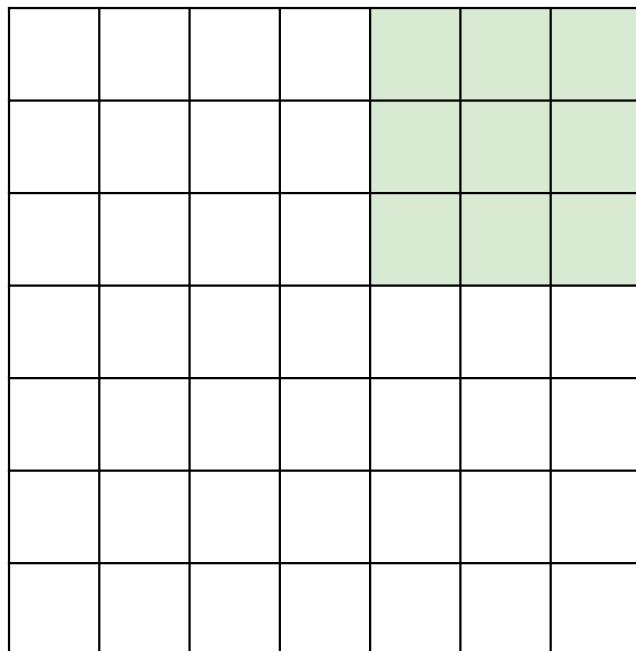
7x7 input
assume 3x3 connectivity, stride 1

Replicate this column of hidden neurons across space, with some **stride**.



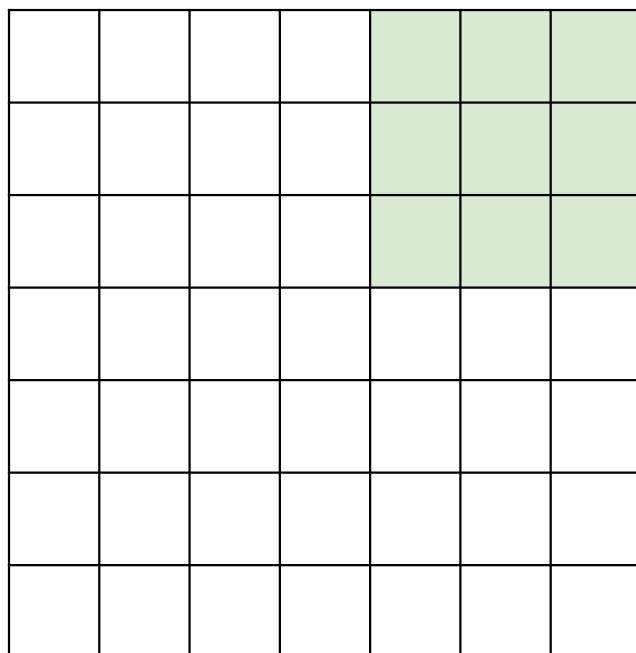
7x7 input
assume 3x3 connectivity, stride 1

Replicate this column of hidden neurons across space, with some **stride**.



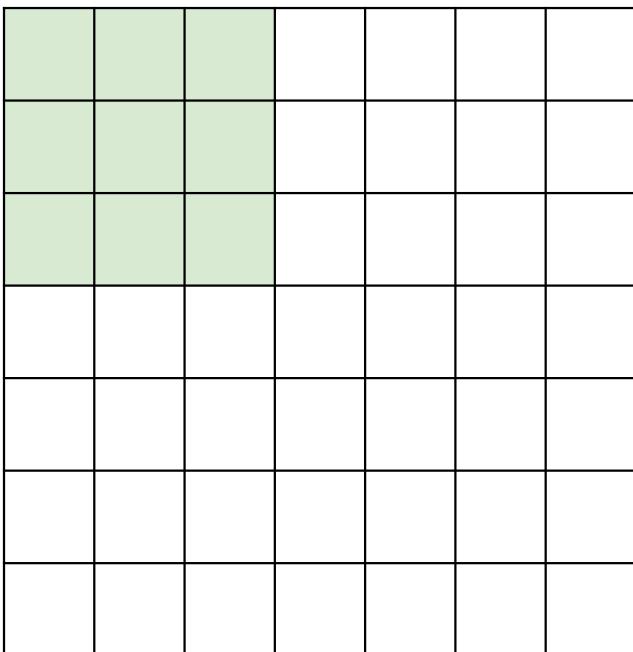
7x7 input
assume 3x3 connectivity, stride 1

Replicate this column of hidden neurons across space, with some **stride**.



7x7 input
assume 3x3 connectivity, stride 1
=> **5x5 output**

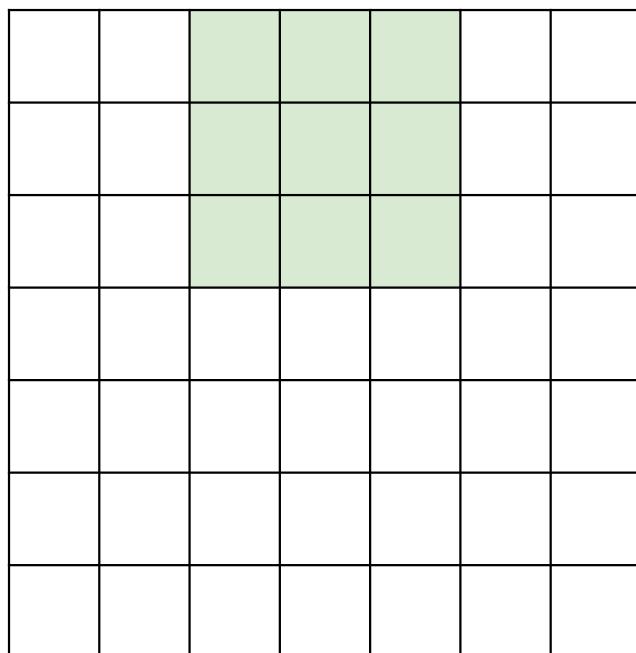
Replicate this column of hidden neurons across space, with some **stride**.



7x7 input
assume 3x3 connectivity, stride 1
=> **5x5 output**

what about stride 2?

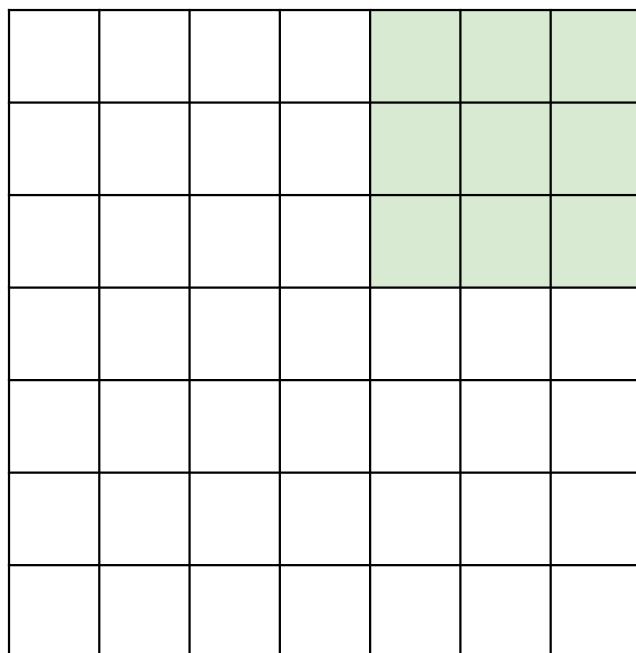
Replicate this column of hidden neurons across space, with some **stride**.



7x7 input
assume 3x3 connectivity, stride 1
=> **5x5 output**

what about stride 2?

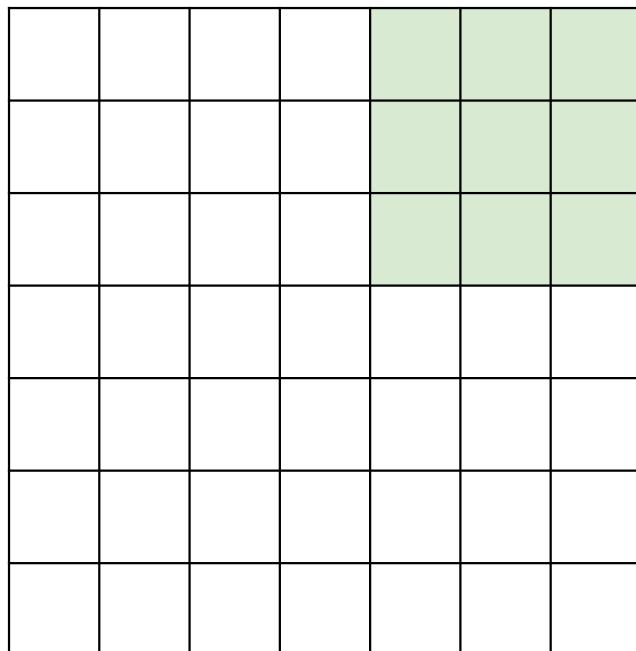
Replicate this column of hidden neurons across space, with some **stride**.



7x7 input
assume 3x3 connectivity, stride 1
=> **5x5 output**

what about stride 2?

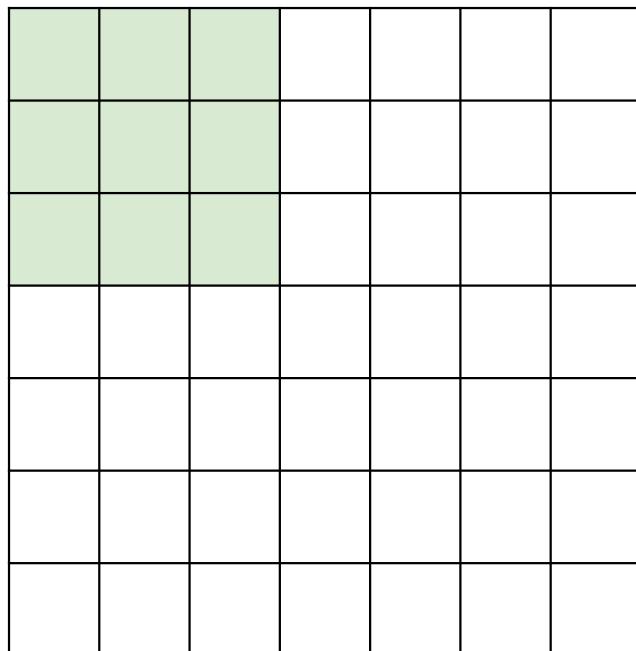
Replicate this column of hidden neurons across space, with some **stride**.



7x7 input
assume 3x3 connectivity, stride 1
=> **5x5 output**

what about stride 2?
=> **3x3 output**

Replicate this column of hidden neurons across space, with some **stride**.

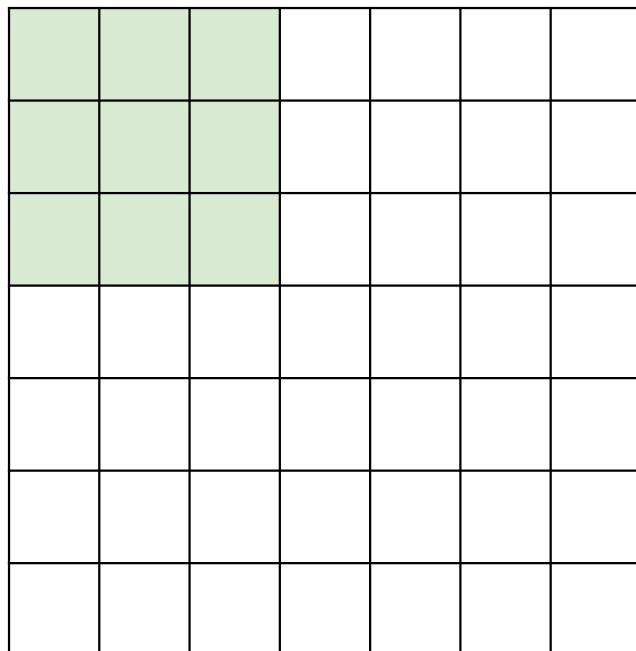


7x7 input
assume 3x3 connectivity, stride 1
=> **5x5 output**

what about stride 2?
=> **3x3 output**

what about stride 3?

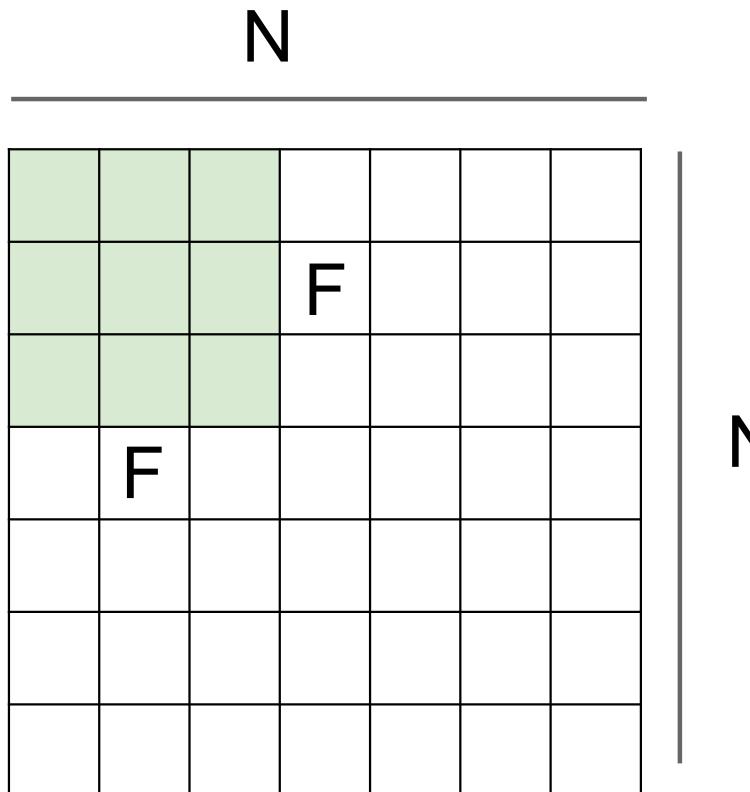
Replicate this column of hidden neurons across space, with some **stride**.



7x7 input
assume 3x3 connectivity, stride 1
=> **5x5 output**

what about stride 2?
=> **3x3 output**

what about stride 3? **Cannot.**



Output size:
 $(N - F) / \text{stride} + 1$

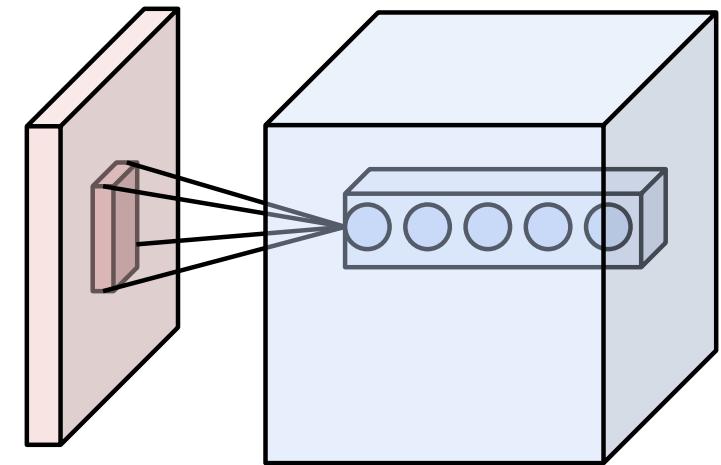
e.g. $N = 7$, $F = 3$:
stride 1 => $(7 - 3)/1 + 1 = 5$
stride 2 => $(7 - 3)/2 + 1 = 3$
stride 3 => $(7 - 3)/3 + 1 = \dots$:\

Examples time:

Input volume: **32x32x3**

Receptive fields: **5x5, stride 1**

Number of neurons: **5**



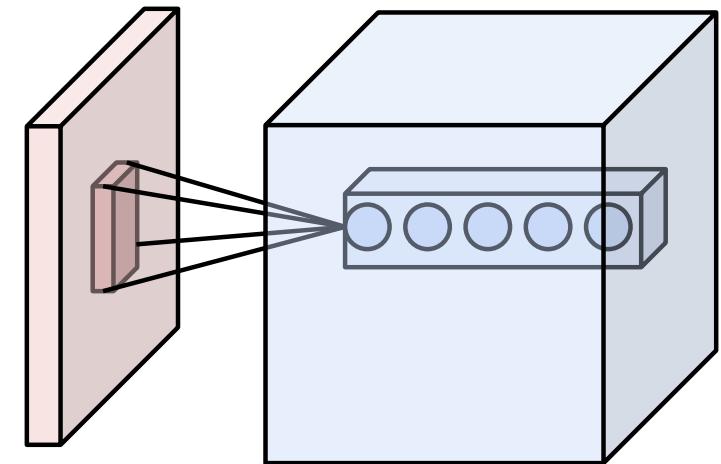
Output volume: ?

Examples time:

Input volume: **32x32x3**

Receptive fields: **5x5, stride 1**

Number of neurons: **5**



Output volume: $(32 - 5) / 1 + 1 = 28$, so: **28x28x5**

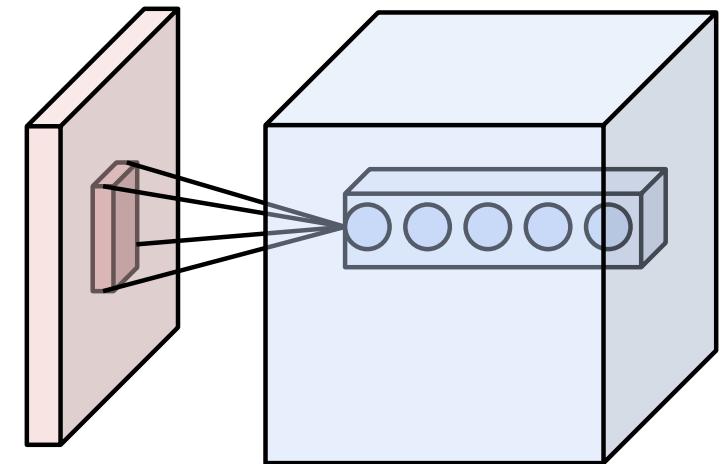
How many weights for each of the 28x28x5 neurons?

Examples time:

Input volume: **32x32x3**

Receptive fields: **5x5, stride 1**

Number of neurons: **5**



Output volume: $(32 - 5) / 1 + 1 = 28$, so: **28x28x5**

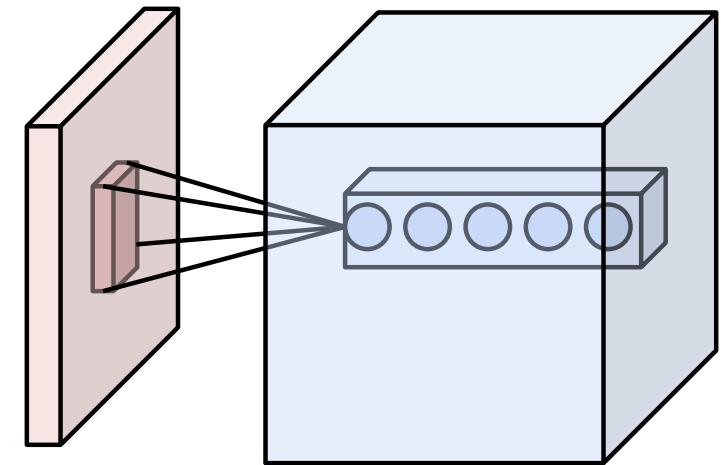
How many weights for each of the 28x28x5
neurons? **5x5x3 = 75**

Examples time:

Input volume: **32x32x3**

Receptive fields: **5x5, stride 2**

Number of neurons: **5**



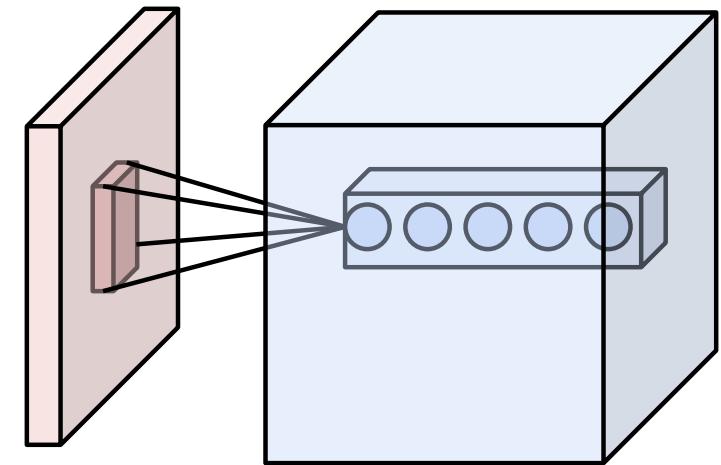
Output volume: ?

Examples time:

Input volume: **32x32x3**

Receptive fields: **5x5, stride 2**

Number of neurons: **5**



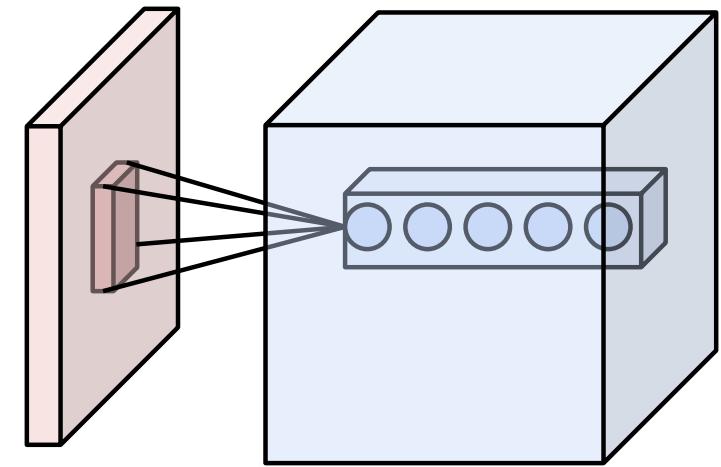
Output volume: ? **Cannot**: $(32-5)/2 + 1 = 14.5$:\

Examples time:

Input volume: **32x32x3**

Receptive fields: **5x5, stride 3**

Number of neurons: **5**



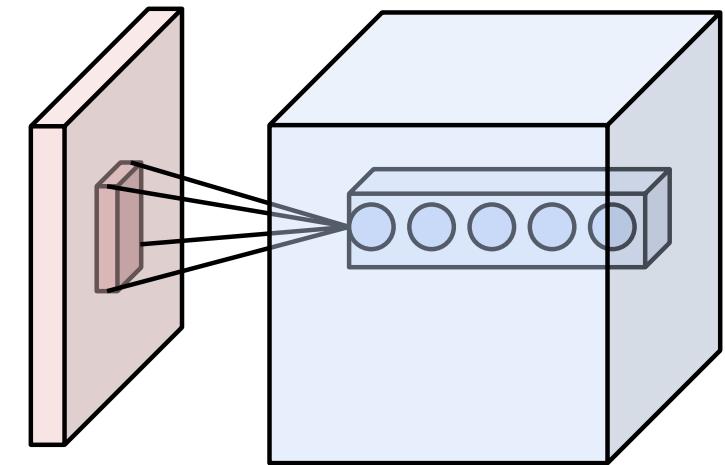
Output volume: ?

Examples time:

Input volume: **32x32x3**

Receptive fields: **5x5, stride 3**

Number of neurons: **5**



Output volume: $(32 - 5) / 3 + 1 = 10$, so: **10x10x5**

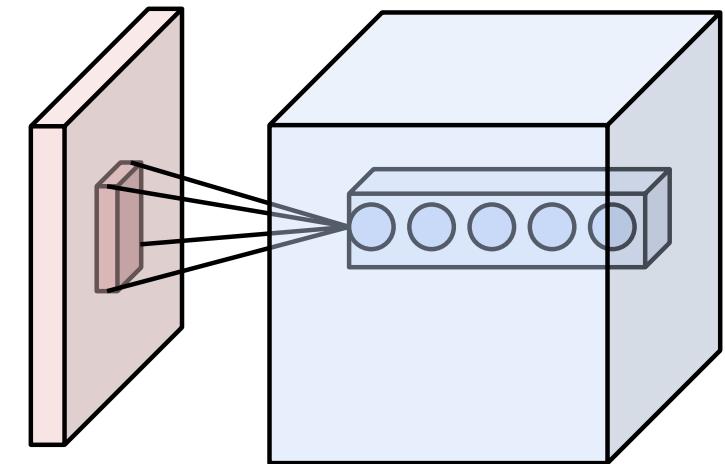
How many weights for each of the 10x10x5 neurons?

Examples time:

Input volume: **32x32x3**

Receptive fields: **5x5, stride 3**

Number of neurons: **5**



Output volume: $(32 - 5) / 3 + 1 = 10$, so: **10x10x5**

How many weights for each of the 10x10x5 neurons? **5x5x3 = 75** (unchanged)

In practice: Common to zero pad the border

0	0	0	0	0	0			
0								
0								
0								
0								

(in each channel)

e.g. input 7x7

neuron with receptive field 3x3, stride 1

pad with 1 pixel border => what is the output?

In practice: Common to zero pad the border

0	0	0	0	0	0			
0								
0								
0								
0								

(in each channel)

e.g. input 7x7

neuron with receptive field 3x3, stride 1
pad with 1 pixel border => what is the output?

7x7 => preserved size!

in general, common to see stride 1, size F, and
zero-padding with $(F-1)/2$.
(Will preserve input size spatially)

“Same convolution” (preserves size)

Input [9x9]

3x3 neurons, stride 1, pad **1** =>

[9x9]

3x3 neurons, stride 1, pad **1** =>

[9x9]

- No headaches when sizing architectures
- Works well

“Valid convolution” (shrinks size)

Input [9x9]

3x3 neurons, stride 1, pad **0** =>

[7x7]

3x3 neurons, stride 1, pad **0** =>

[5x5]

- **Headaches** with sizing the full architecture
- **Works Worse!** Border information will “wash away”, since those values are only used once in the forward function

Summary:

Input volume of size $[W_1 \times H_1 \times D_1]$

using K neurons with receptive fields $F \times F$ and applying them at strides of S gives

Output volume: $[W_2, H_2, D_2]$

$$W_2 = (W_1 - F)/S + 1$$

$$H_2 = (H_1 - F)/S + 1$$

$$D_2 = K$$

There's one more problem...

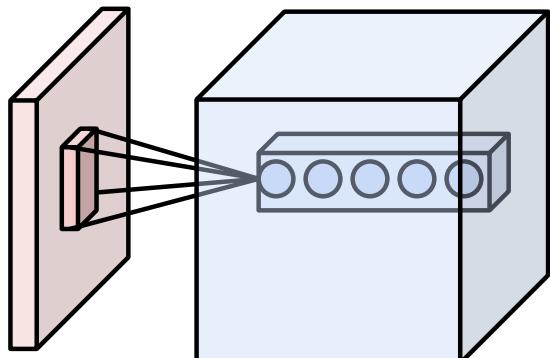
Assume input $[32 \times 32 \times 3]$

30 neurons with receptive fields **5x5**, applied at **stride 1/pad1**:

=> Output volume: $[32 \times 32 \times 30]$ ($32 \times 32 \times 30 = 30720$ neurons)

Each neuron has $5 \times 5 \times 3$ (=75) weights

=> Number of weights in such layer: $30720 \times 75 \approx 3 \text{ million :}$



There's one more problem...

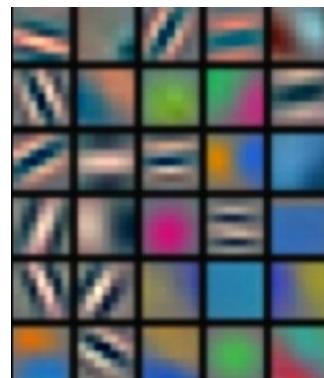
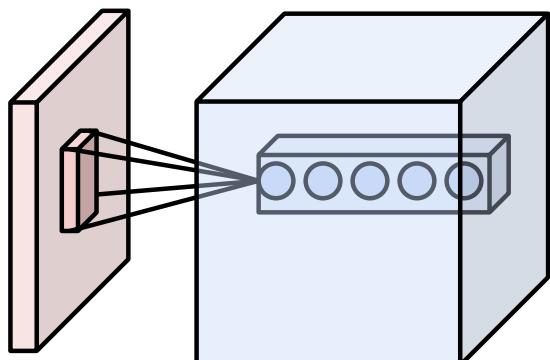
Assume input $[32 \times 32 \times 3]$

30 neurons with receptive fields 5×5 , applied at **stride 1/pad 2**:

=> Output volume: $[32 \times 32 \times 30]$ ($32 \times 32 \times 30 = 30720$ neurons)

Each neuron has $5 \times 5 \times 3$ (=75) weights

=> Number of weights in such layer: $30720 \times 75 \approx 3 \text{ million :)}$



Example trained filters

There's one more problem...

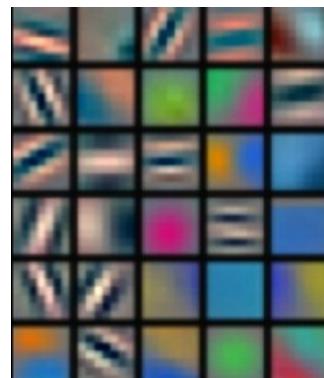
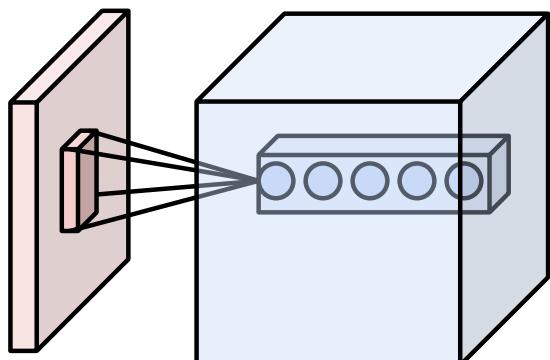
Assume input $[32 \times 32 \times 3]$

30 neurons with receptive fields **5x5**, applied at **stride 1/pad1**:

=> Output volume: $[32 \times 32 \times 30]$ ($32 \times 32 \times 30 = 30720$ neurons)

Each neuron has $5 \times 5 \times 3$ (=75) weights

=> Number of weights in such layer: $30720 \times 75 \approx 3$ million :\\



← Example trained weights
IDEA: let's not learn the same thing across all spatial locations

Our first ConvNet layer had size **[32 x 32 x3]**

If we had **30** neurons with receptive fields **5x5, stride 1, pad 1**

Output volume: **[32 x 32 x 30]** ($32*32*30 = 30720$ neurons)

Each neuron has **5*5*3 (=75)** weights

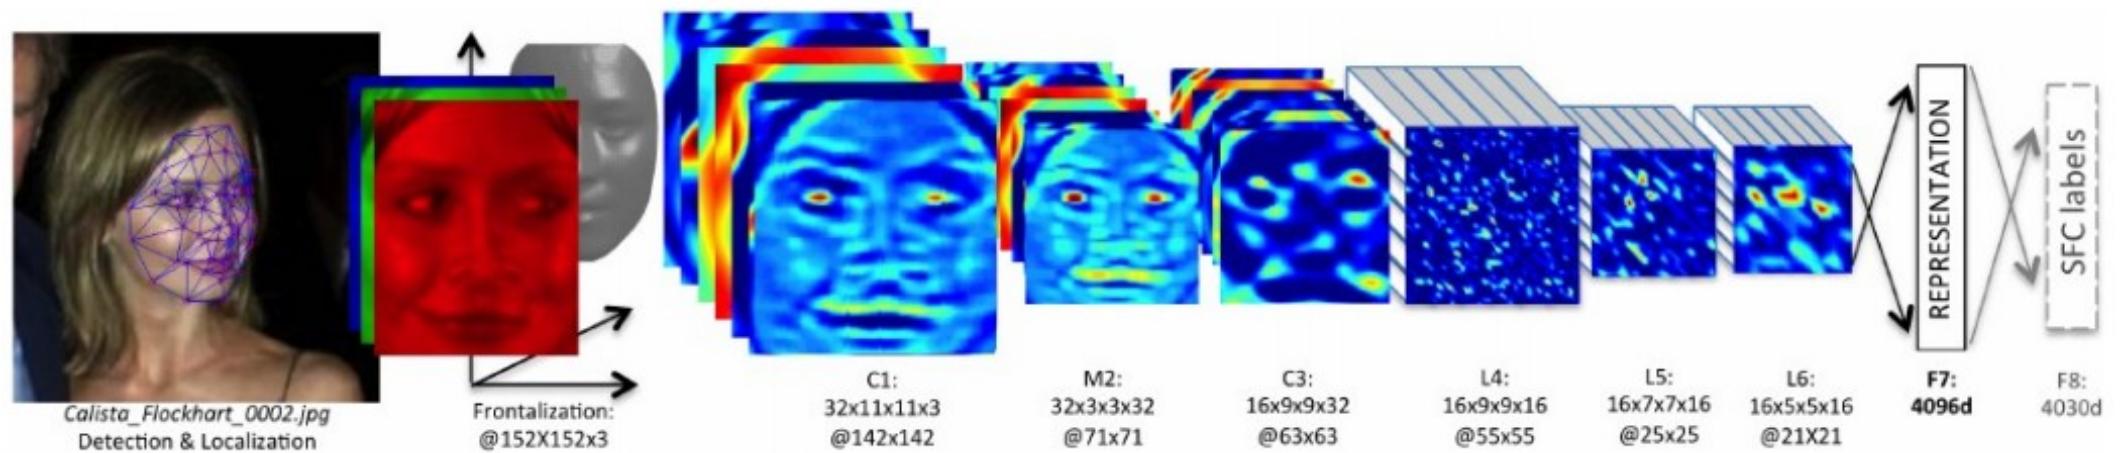
Before:

#weights in such layer: **(32*32*30) * 75 = 3 million** :\

Now: (paramater sharing)

#weights in the layer: **30 * 75 = 2250.**

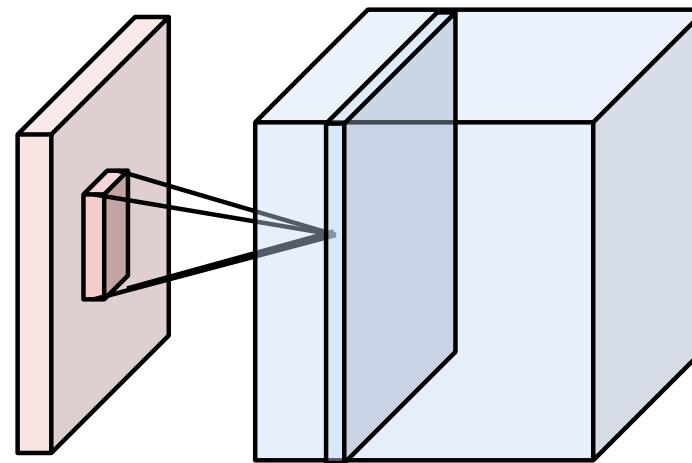
Note: sometimes it's not a good idea to share the parameters



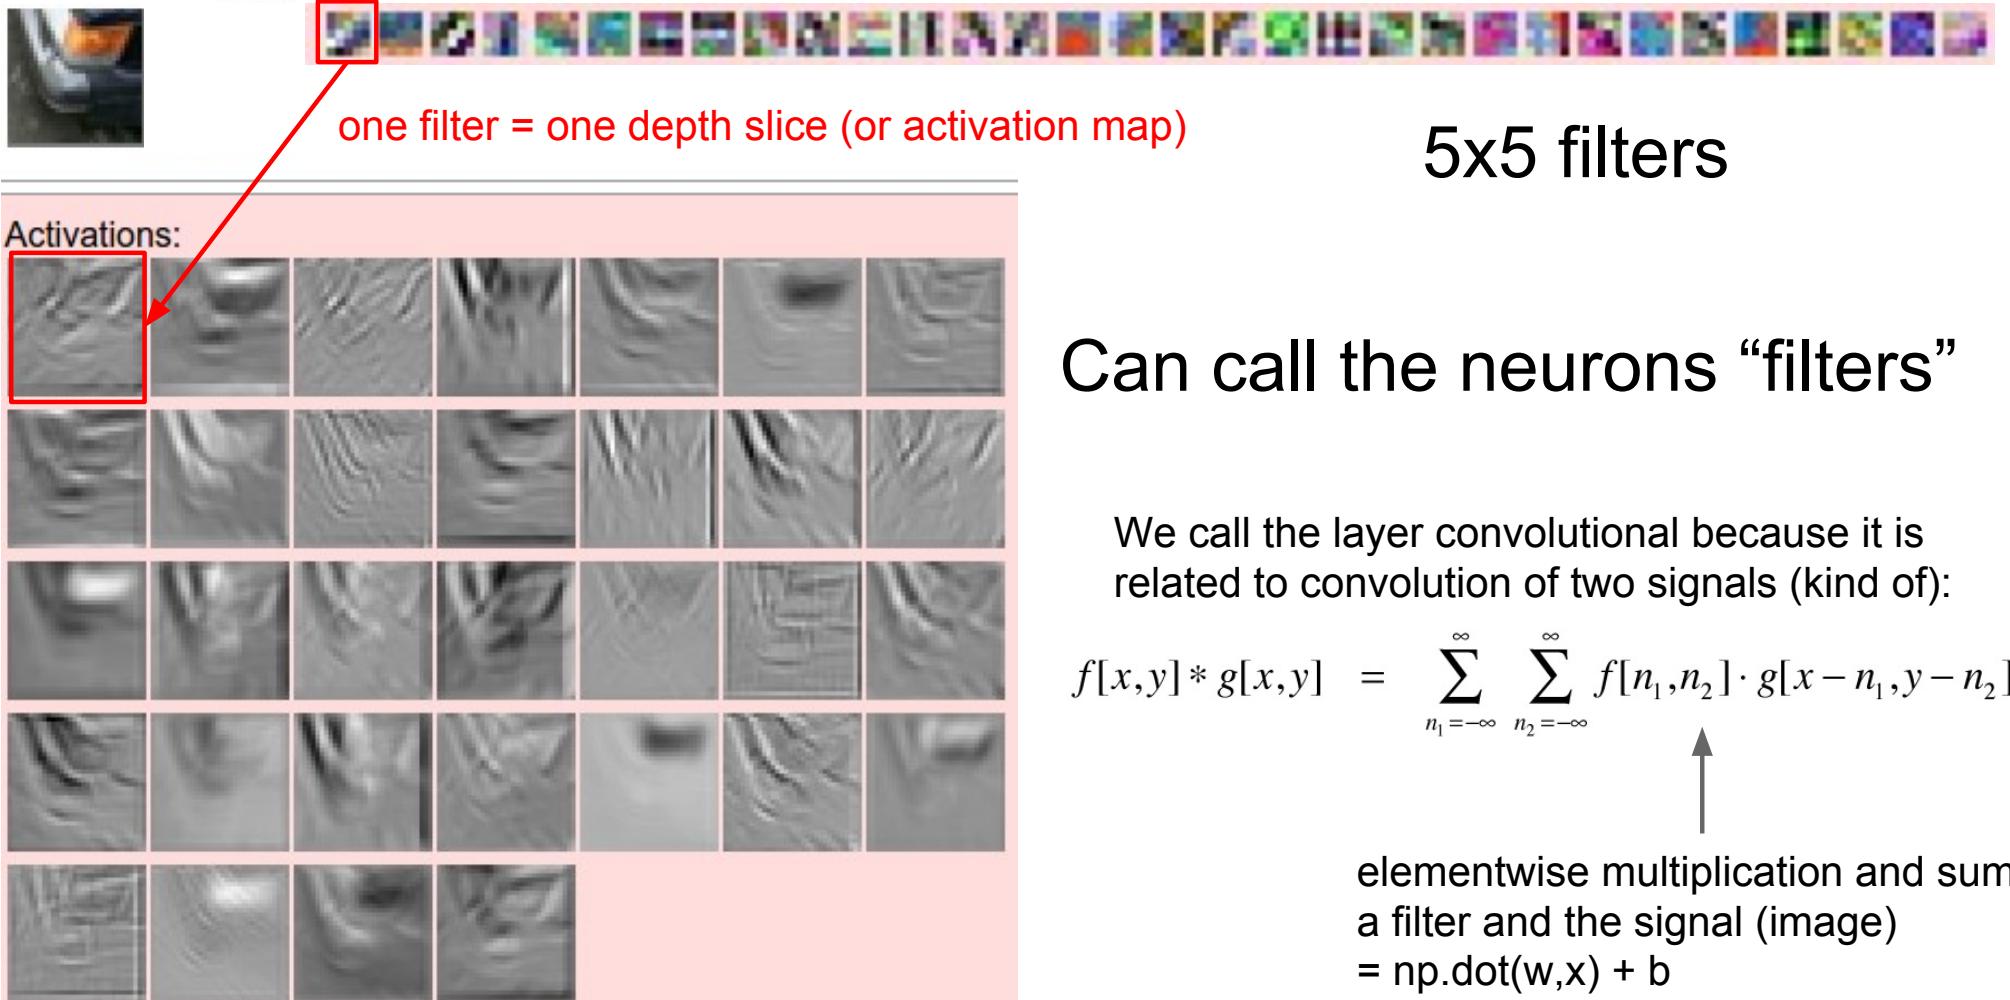
We'd like to be able to learn different things at different spatial positions

These layers are called **Convolutional Layers**

1. Connect neurons only to local receptive fields
2. Use the same neuron weight parameters for neurons in each “depth slice” (i.e. across spatial positions)

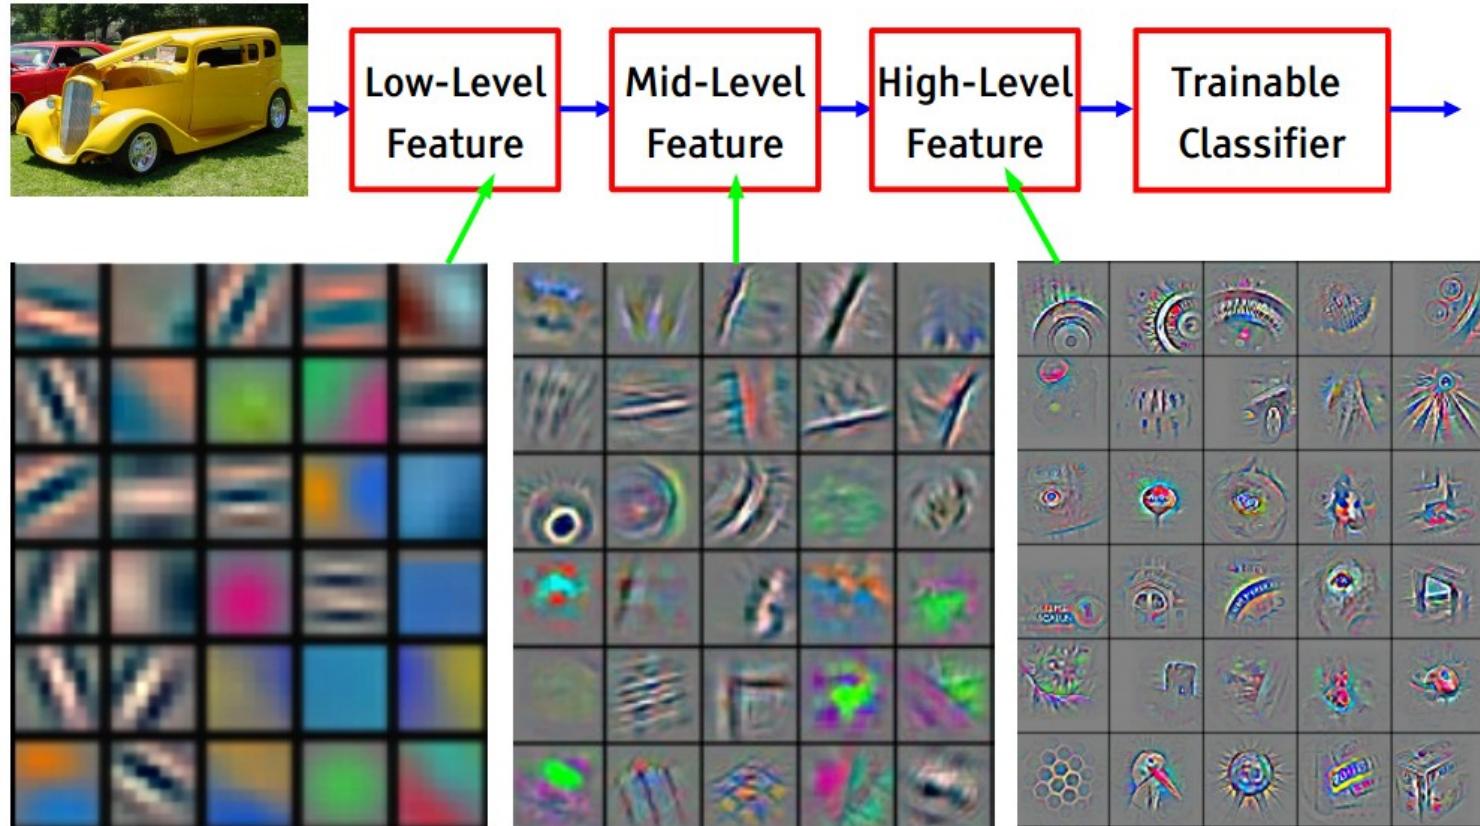


one activation map (a depth slice),
computed with one set of weights



Fast-forward to today

[From recent Yann LeCun slides]

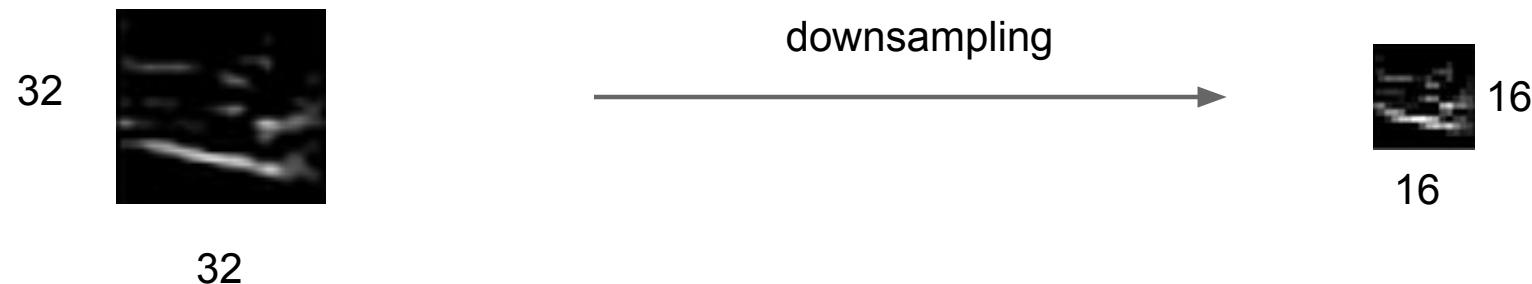


Feature visualization of convolutional net trained on ImageNet from [Zeiler & Fergus 2013]



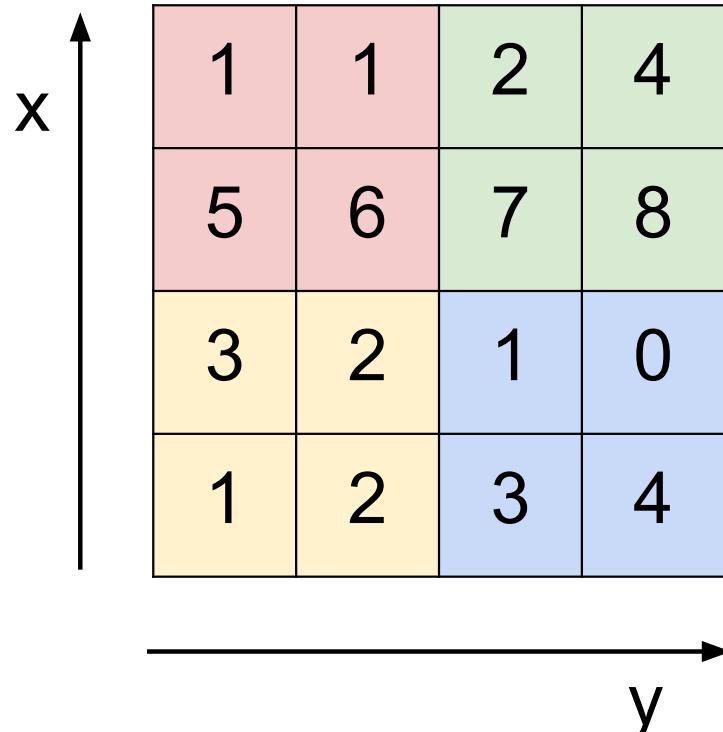
In ConvNet architectures, **Conv** layers are often followed by **Pool** layers

- convenience layer: makes the representations smaller and more manageable without losing too much information. Computes MAX operation (most common)

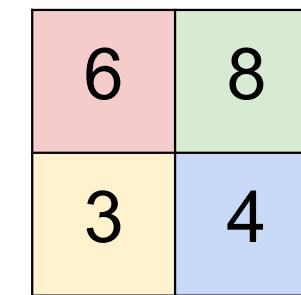


MAX POOLING

Single depth slice



max pool with 2x2 filters
and stride 2



In ConvNet architectures, **Conv** layers are often followed by **Pool** layers

- convenience layer: makes the representations smaller and more manageable without losing too much information. Computes MAX operation (most common)

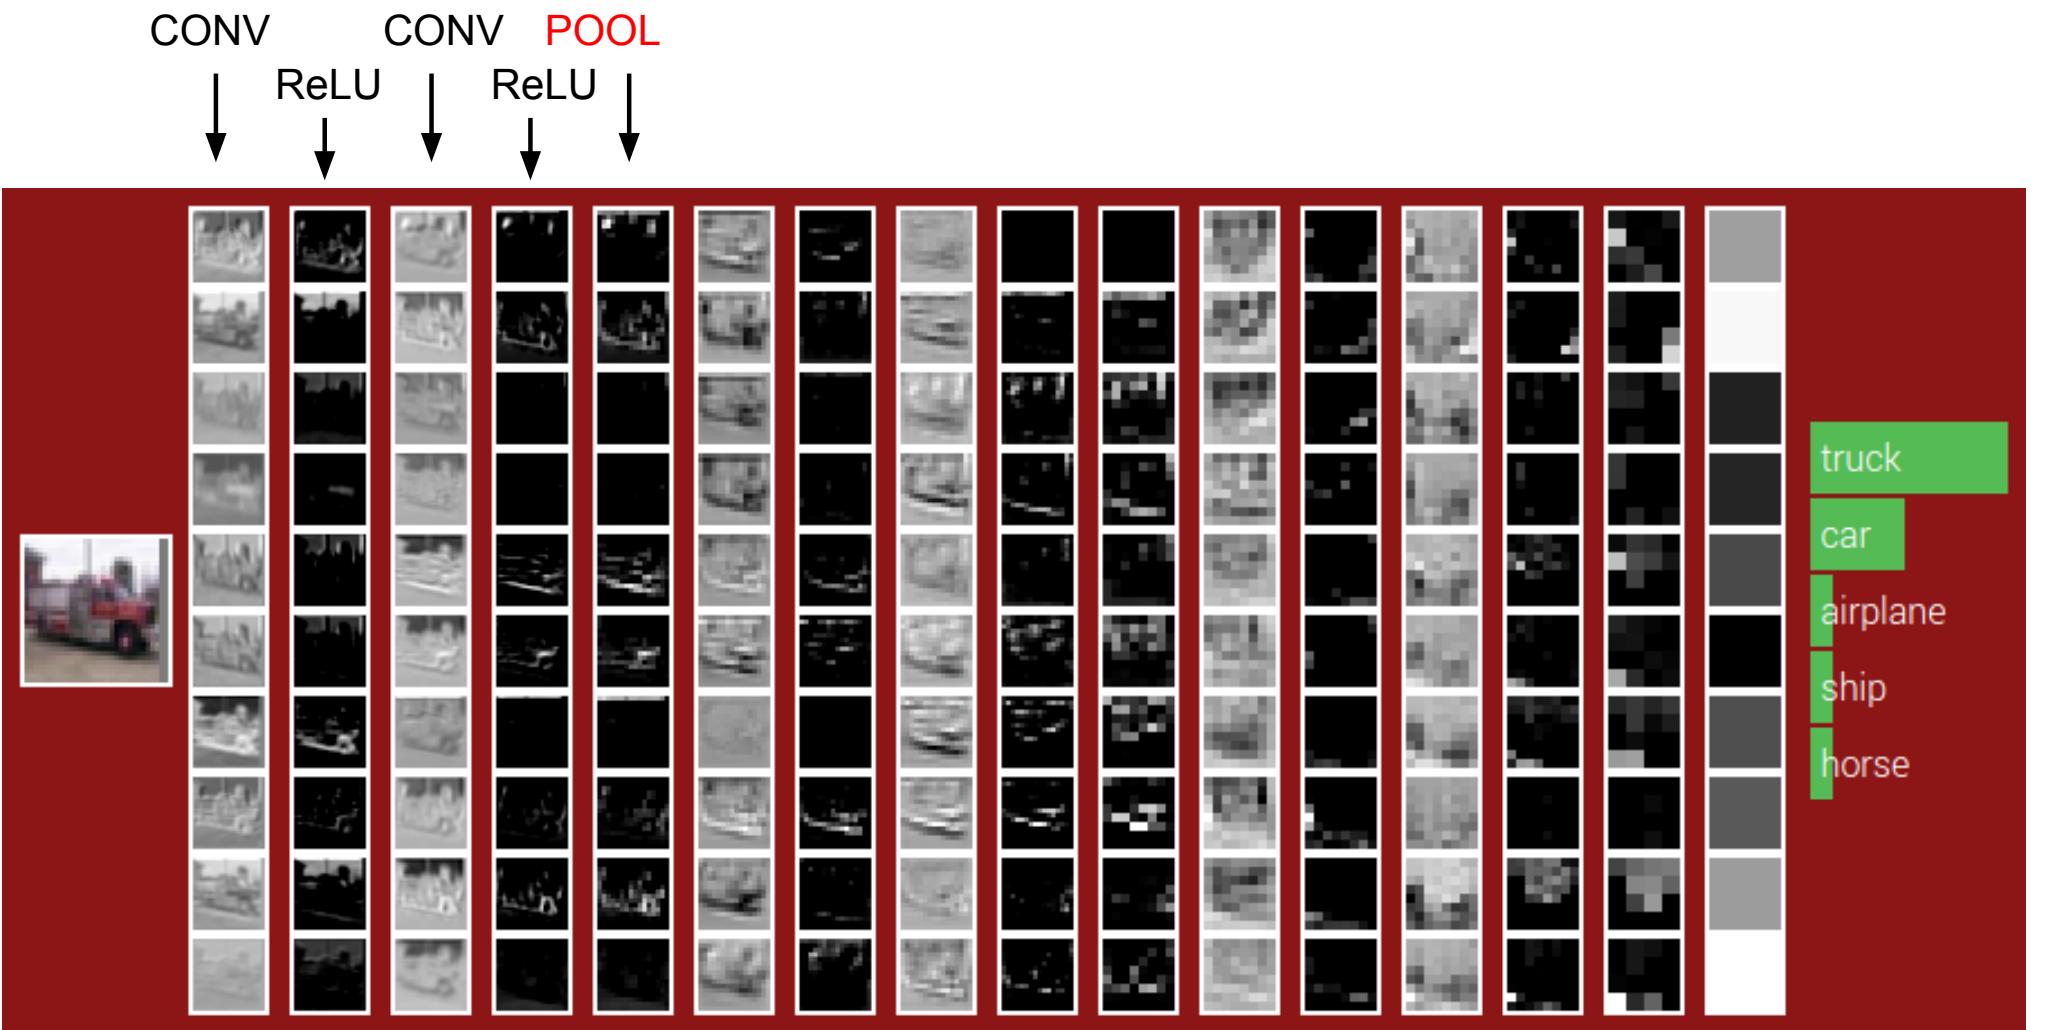
Input volume of size [W1 x H1 x D1]

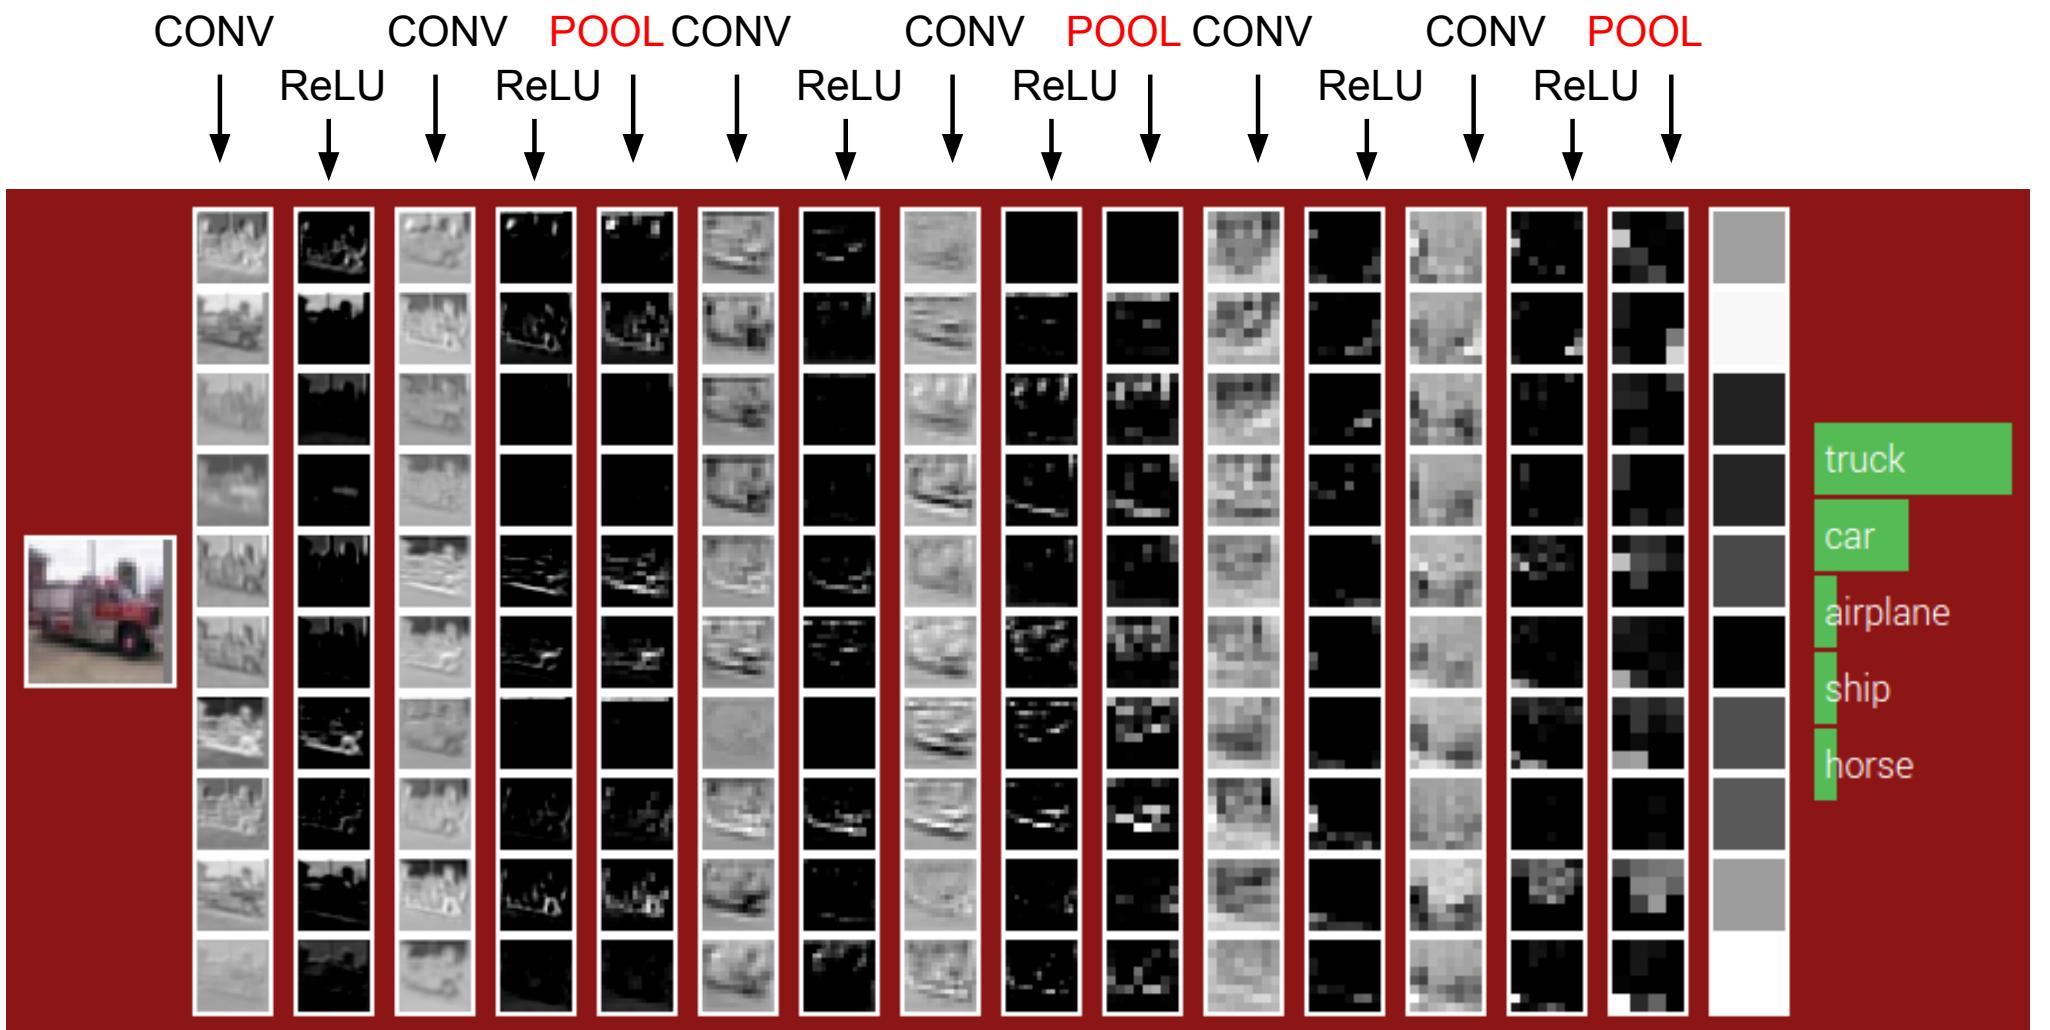
Pooling unit receptive fields $F \times F$ and applying them at strides of S gives

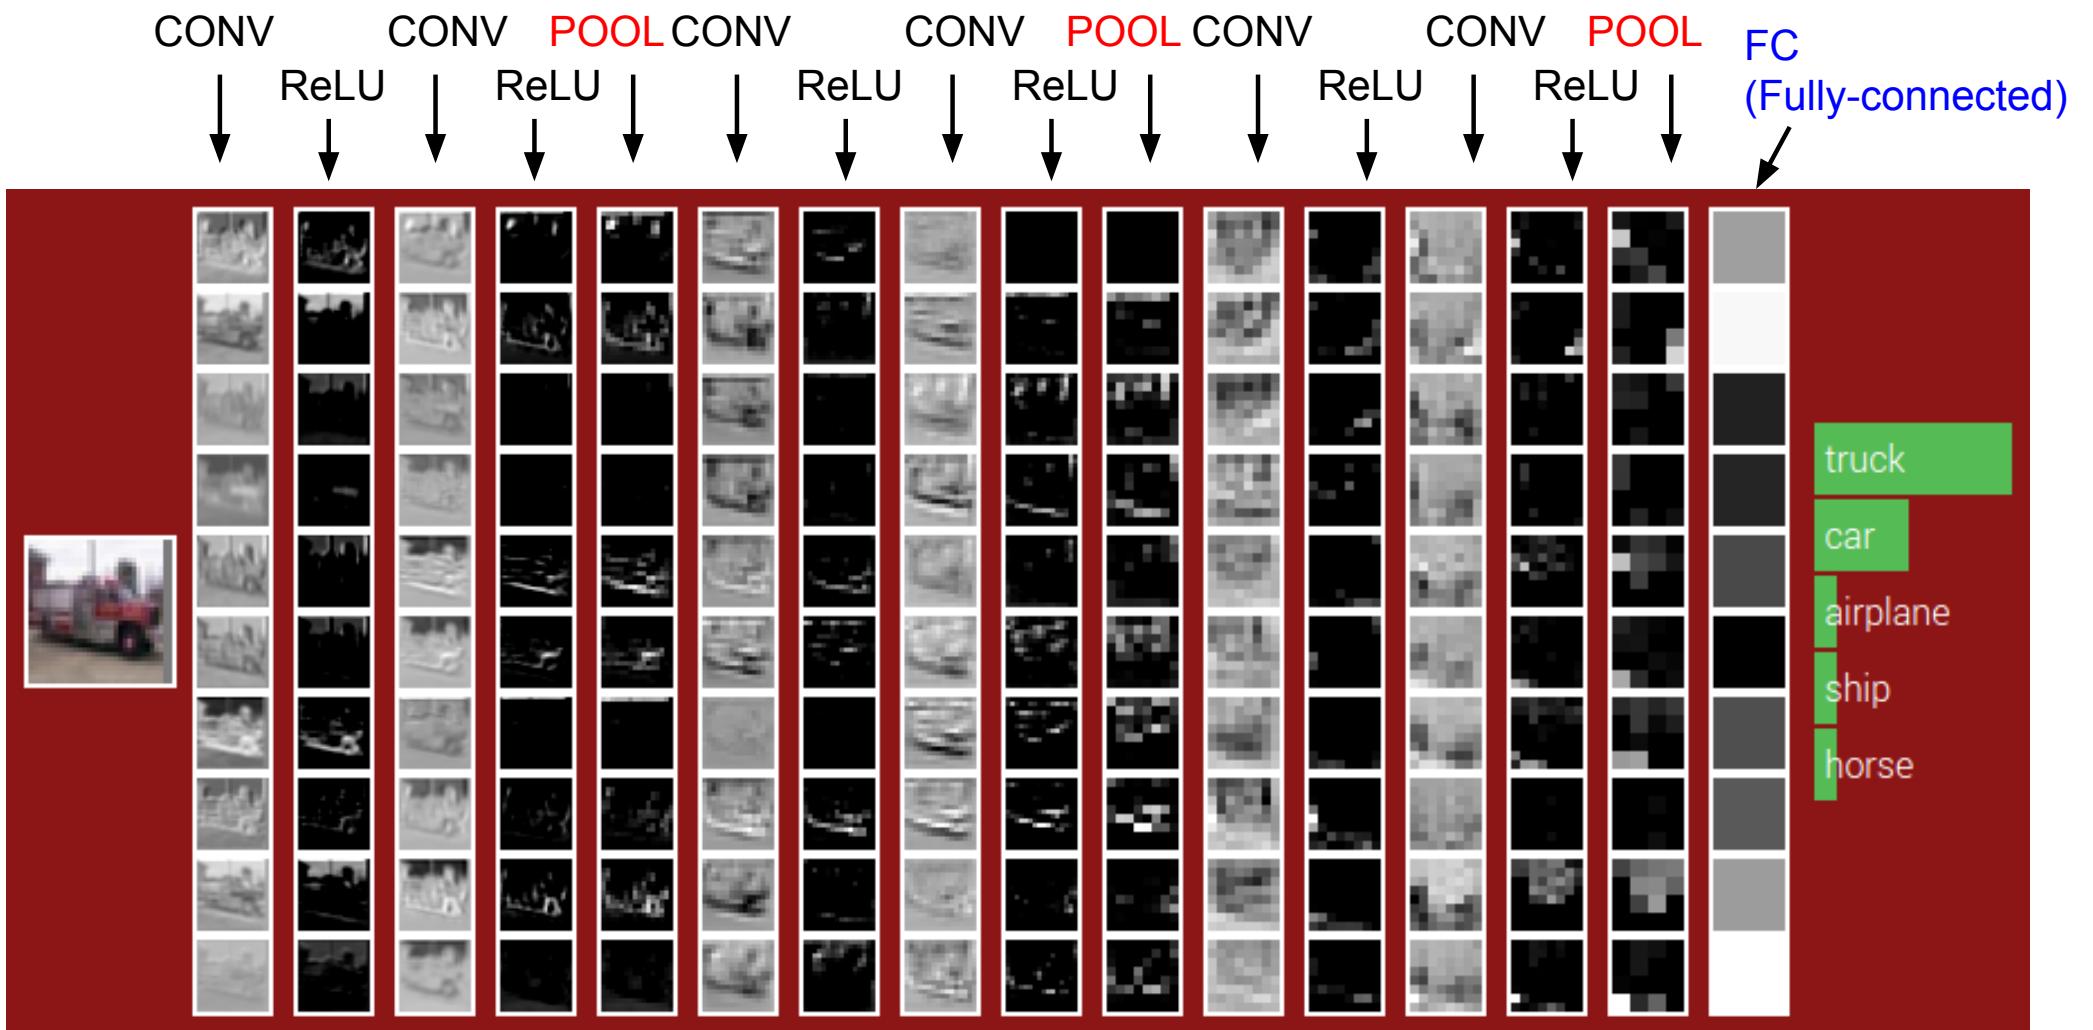
Output volume: [W2, H2, D1]

$$W2 = (W1 - F)/S + 1, H2 = (H1 - F)/S + 1$$

Note: pooling happens independently across each slice, preserving number of slices
E.g. a pooling “neuron” of size 2x2 will perform MAX operation over 4 numbers.







Modern CNNs:

- use **filter sizes of 3x3** (maybe even 2x2 or 1x1!)
- use **pooling sizes of 2x2** (maybe even less - e.g. fractional pooling!)
- **stride 1**
- **very deep**

(if too expensive for time/space, might have to downsample more, or make bigger strides, etc)

Eliminate sizing headaches TIPS/TRICKS

- start with image that has power-of-2 size
- for **conv layers**, use stride 1 filter size 3x3 pad input with a border of zeros (1 spatially)

This makes it so that: [W1,H1,D1] -> [W1,H1,D2] (i.e. spatial size exactly preserved)

- for **pool layers**, use pool size 2x2 (more = worse)

For big images, good heuristic is to violate the previous rules but only for the first layer. E.g. “AlexNet”:

input 227x227x3

first CONV layer: 96 11x11 filters, stride 4

$$\Rightarrow (227 - 11)/4 + 1 = 55$$

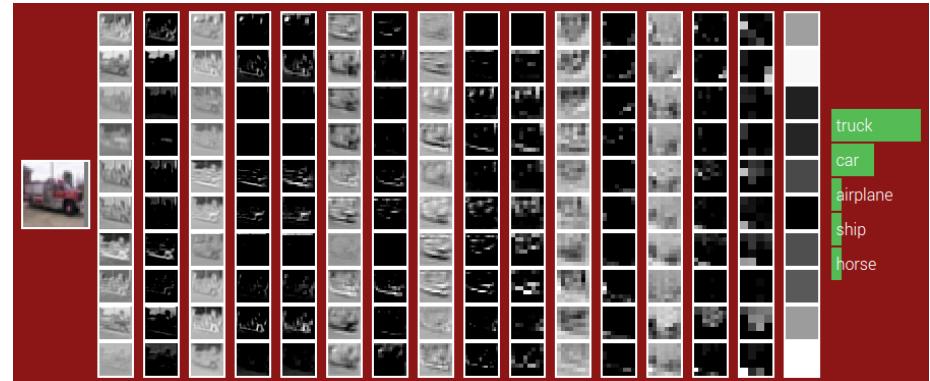
\Rightarrow output volume 55 x 55 x 96

...and from then stride 1 “same” convolutions (with padding).

Example:

input: [32x32x3]

CONV with 10 3x3 filters, stride 1, pad 1:
gives:



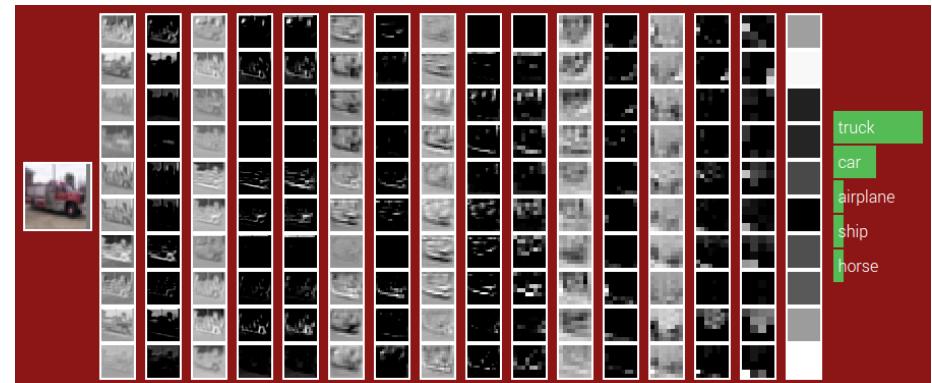
Example:

input: [32x32x3]

CONV with 10 3x3 filters, stride 1, pad 1:

gives: [32x32x10]

new parameters:



Example:

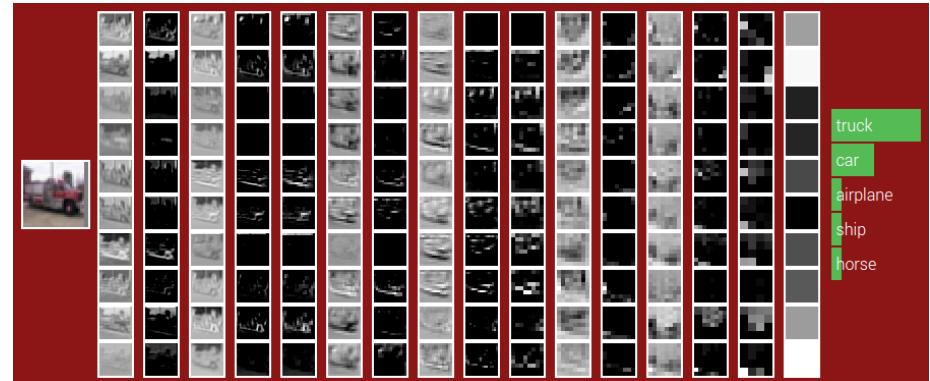
input: [32x32x3]

CONV with 10 3x3 filters, stride 1, pad 1:

gives: [32x32x10]

new parameters: $(3 \times 3 \times 3) \times 10 + 10 = 280$

RELU



Example:

input: [32x32x3]

CONV with 10 3x3 filters, stride 1, pad 1:

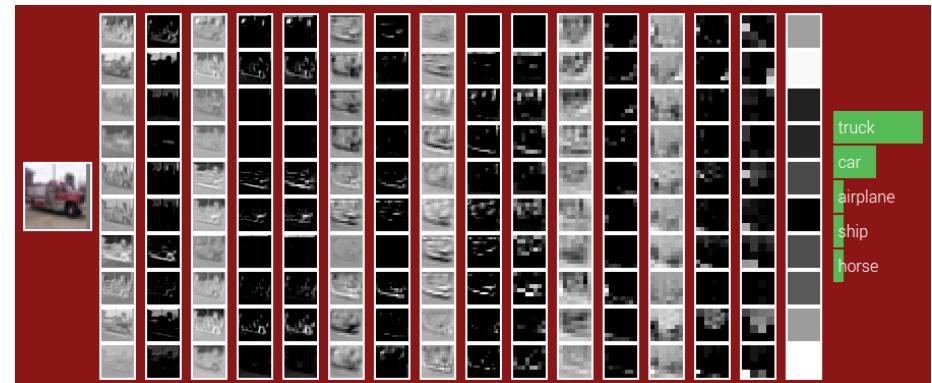
gives: [32x32x10]

new parameters: $(3 \times 3 \times 3) \times 10 + 10 = 280$

RELU

CONV with 10 3x3 filters, stride 1, pad 1:

gives:



Example:

input: [32x32x3]

CONV with 10 3x3 filters, stride 1, pad 1:

gives: [32x32x10]

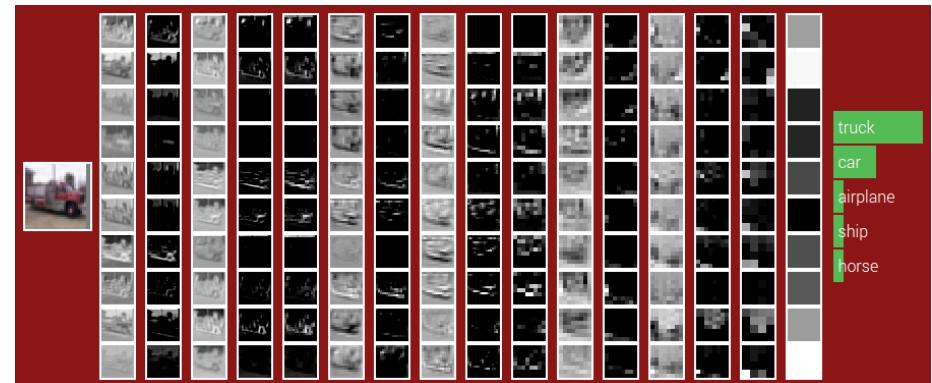
new parameters: $(3 \times 3 \times 3) \times 10 + 10 = 280$

RELU

CONV with 10 3x3 filters, stride 1, pad 1:

gives: [32x32x10]

new parameters:



Example:

input: [32x32x3]

CONV with 10 3x3 filters, stride 1, pad 1:

gives: [32x32x10]

new parameters: $(3 \times 3 \times 3) \times 10 + 10 = 280$

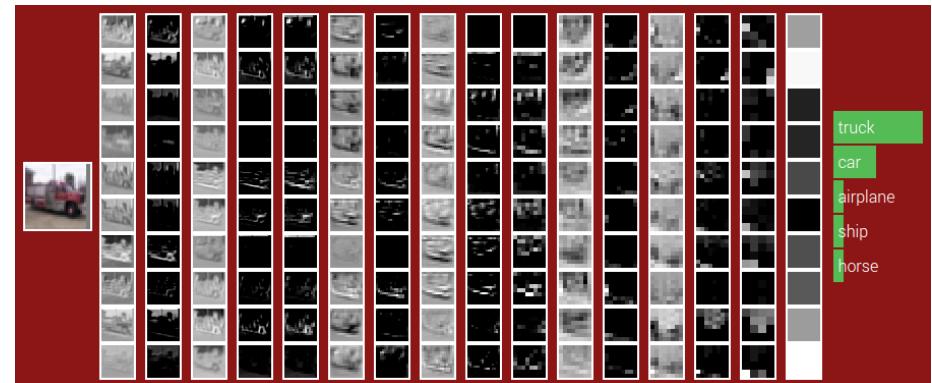
RELU

CONV with 10 3x3 filters, stride 1, pad 1:

gives: [32x32x10]

new parameters: $(3 \times 3 \times 10) \times 10 + 10 = 910$

RELU



Example:

input: [32x32x3]

CONV with 10 3x3 filters, stride 1, pad 1:

gives: [32x32x10]

new parameters: $(3 \times 3 \times 3) \times 10 + 10 = 280$

RELU

CONV with 10 3x3 filters, stride 1, pad 1:

gives: [32x32x10]

new parameters: $(3 \times 3 \times 10) \times 10 + 10 = 910$

RELU

POOL with 2x2 filters, stride 2:

gives:

